



```
1 install.packages("e1071")

Installing package into ‘/usr/local/lib/R/site-library’
(as ‘lib’ is unspecified)

also installing the dependency ‘proxy’



1 library(e1071)
2 set.seed (1)
3 # We now use the svm() function to fit the support vector classifier for a given value of the cost parameter.
4 # Here we demonstrate the use of this function on a two-dimensional example so that we can plot the resulting
5 # decision boundary.
6 # We begin by generating the observations, which belong to two classes.
7 x=matrix(rnorm(20*2), ncol=2)
8 y=c(rep(-1,10), rep(1,10))
9 x[y==1,]=x[y==1,] + 1
10 x
11 y
12 # We begin by checking whether the classes are linearly separable.
13 plot(x, col=(3-y))
14 # They are not. Next, we fit the support vector classifier.
15 # Note that in order for the svm() function to perform classification
16 # we must encode the response as a factor variable.
17 # We now create a data frame with the response coded as a factor.
18 dat <- data.frame(x = x,y = as.factor(y))
19 svmfit <- svm(y ~., data=dat, kernel="linear", cost=10,scale=FALSE)
20 # The argument scale=FALSE tells the svm() function not to scale each feature to
21 # have mean zero or standard deviation one;
22 # depending on the application, one might prefer to use scale=TRUE.
23 # We can now plot the support vector classifier obtained:
24 plot(svmfit , dat)
25 # Note that the two arguments to the plot.svm() function are the output of the call to svm(),
26 #as well as the data used in the call to svm().
27 # The region of feature space that will be assigned to the -1 class is shown in light blue,
28 # and the region that will be assigned to the +1 class is shown in purple.
```

0.5757814 -0.47815006
-0.3053884 0.41794156
2.5117812 2.35867955
1.3898432 0.89721227
0.3787594 1.38767161
-1.2146999 0.94619496
2.1249309 -0.37705956
0.9550664 0.58500544
0.9838097 0.60571005
1.9438362 0.94068660
1.8212212 2.10002537
1.5939013 1.76317575
-1 · -1 · -1 · -1 · -1 · -1 · -1 · -1 · -1 · -1 · 1 · 1 · 1 · 1 · 1 · 1 · 1 · 1 · 1 · 1



```
1 svmfit <- svm(y ~., data=dat, kernel="linear", cost = 0.1, scale=FALSE)
2 plot(svmfit , dat)
3 svmfit$index
```

1 · 2 · 3 · 4 · 5 · 7 · 9 · 10 · 12 · 13 · 14 · 15 · 16 · 17 · 18 · 20



```
1 set.seed (1)
2 tune.out <- tune(svm, y ~.,data=dat,kernel="linear", ranges=list(cost=c(0.001, 0.01, 0.1, 1,5,10,100)))
3 # We can easily access the cross-validation errors for each of these models using the summary() command:
4 summary(tune.out)
5 # We see that cost=0.1 results in the lowest cross-validation error rate.
6 # The tune() function stores the best model obtained, which can be accessed as follows:
7 bestmod=tune.out$best.model
8 summary(bestmod)
9 # The predict() function can be used to predict the class label on a set of test observations,
10 # at any given value of the cost parameter. We begin by generating a test data set.
11 xtest=matrix(rnorm(20*2), ncol=2)
12 ytest=sample(c(-1,1), 20, rep=TRUE)
13 xtest[ytest==1,]=xtest[ytest==1,] + 1
14 testdat=data.frame(x=xtest, y=as.factor(ytest))
15 # Now we predict the class labels of these test observations.
16 # Here we use the best model obtained through cross-validation in order to make predictions.
17 ypred <-predict(bestmod ,testdat)
18 table(predict=ypred, truth=testdat$y)
19 #Thus, with this value of cost, 19 of the test observations are correctly classified.
20 # What if we had instead used cost= 0.01?
21 svmfit <- svm(y~., data=dat, kernel="linear", cost=.01, scale=FALSE)
22 ypred=predict(svmfit ,testdat)
23 table(predict=ypred, truth=testdat$y)
24 # In this case one additional observation is misclassified.
25 # Now consider a situation in which the two classes are linearly separable.
26 # Then we can find a separating hyperplane using the svm() function.
27 # We first further separate the two classes in our simulated data so that they are linearly separable:
28 x[y==1,]=x[y==1,]+0.5
29 plot(x, col=(y+5)/2, pch=19)
```

```
- sampling method: 10-fold cross validation

- best parameters:
cost
0.1

1 #Now the observations are just barely linearly separable.
2 # We fit the support vector classifier and plot the resulting hyperplane,
3 # using a very large value of cost so that no observations are misclassified.
4 dat=data.frame(x=x,y=as.factor(y))
5 svmfit <-svm(y~., data=dat, kernel="linear", cost=1e5)
6 summary(svmfit)
7 plot(svmfit,dat)
8 # No training errors were made and only three support vectors were used.
9 # However, we can see from the figure that the margin is
10 # very narrow (because the observations that are not support vectors, indicated as circles, are very
11 # close to the decision boundary). It seems likely that this model will perform poorly on test data.
12 # We now try a smaller value of cost:
13 svmfit <- svm(y~., data=dat, kernel="linear", cost=1)
14 summary(svmfit)
15 plot(svmfit ,dat)
16 # Using cost=1, we misclassify a training observation, but we also obtain a much wider margin and make
17 # use of seven support vectors.
18 # It seems likely that this model will perform better on test data than the model with cost=1e5.
```

Call:
svm(formula = y ~ ., data = dat, kernel = "linear", cost = 1)


Parameters:
SVM-Type: C-classification
SVM-Kernel: linear
cost: 1

Number of Support Vectors: 7

(4 3)

Number of Classes: 2

Levels:
-1 1



```
1 #We now examine the Khan data set, which consists of a number of tissue samples
2 # corresponding to four distinct types of small round blue cell tumors
3 # For each tissue sample, gene expression measurements are available.
4 #The data set consists of training data, xtrain and ytrain, and testing data, xtest and ytest.
5 install.packages("ISLR")
6 library(e1071)
7 library(ISLR)
8 names(Khan)
9 # Let's examine the dimension of the data:
10 # This data set consists of expression measurements for 2,308 genes.
11 # The training and test sets consist of 63 and 20 observations respectively
12 dim(Khan$xtrain )
13 dim(Khan$xtest )
14 length(Khan$ytrain )
15 length(Khan$ytest )
16 table(Khan$ytrain )
17 table(Khan$ytest )
18 # We will use a support vector approach to predict cancer subtype using gene expression measurements.
19 # In this data set, there are a very large number of features relative to the number of observations.
20 # This suggests that we should use a linear kernel, because the additional flexibility that will
21 # result from using a polynomial or radial kernel is unnecessary.
22 dat <- data.frame(x=Khan$xtrain , y = as.factor(Khan$ytrain ))
23 out <- svm(y ~., data=dat, kernel="linear",cost=10)
24 summary(out)
25 # We see that there are no training errors. In fact, this is not surprising, because the large number
26 # of variables relative to the number of observations implies that it is easy to find hyperplanes that
27 # fully separate the classes.
28 # We are most interested not in the support vector classifier's performance on the training observations,
29 # but rather its performance on the test observations.
30 dat.te=data.frame(x=Khan$xtest , y = as.factor(Khan$ytest ))
31 pred.te=predict(out, newdata=dat.te)
32 table(pred.te, dat.te$y)
33 # We see that using cost=10 yields two test set errors on this data

Installing package into ‘/usr/local/lib/R/site-library’
(as ‘lib’ is unspecified)

'xtrain' · 'xtest' · 'ytrain' · 'ytest'
63 · 2308
20 · 2308
63
20

1 2 3 4
8 23 12 20

1 2 3 4
3 6 6 5

Call:
svm(formula = y ~ ., data = dat, kernel = "linear", cost = 10)

Parameters:
  SVM-Type:  C-classification
 SVM-Kernel: linear
      cost: 10

Number of Support Vectors:  58

( 20 20 11 7 )


Number of Classes:  4

Levels:
 1 2 3 4

pred.te 1 2 3 4
1 3 0 0 0
2 0 6 2 0
3 0 0 4 0
4 0 0 0 5

1 n <- 150 # number of data points
2 p <- 2 # dimension
3 sigma <- 1 # variance of the distribution
4 meanpos <- 0 # centre of the distribution of positive examples
5 meanneg <- 3 # centre of the distribution of negative examples
6 npos <- round(n/2) # number of positive examples
7 nneg <- n-npos # number of negative examples
8 # Generate the positive and negative examples
9 xpos <- matrix(rnorm(npos*p,mean=meanpos,sd=sigma),npos,p)
10 xneg <- matrix(rnorm(nneg*p,mean=meanneg,sd=sigma),npos,p)
11 x <- rbind(xpos,xneg)
12 # Generate the labels
13 y <- matrix(c(rep(1,npos),rep(-1,nneg)))
14 # Visualize the data
15 plot(x,col=ifelse(y>0,1,2))
16 legend("topleft",c('Positive', 'Negative'),col=seq(2),pch=1,text.col=seq(2))
17 #
18 ntrain <- round(n*0.8) # number of training examples
19 tindex <- sample(n,ntrain) # indices of training samples
20 xtrain <- x[tindex,]
21 xtest <- x[-tindex,]
22 ytrain <- y[tindex]
23 ytest <- y[-tindex]
24 istrain=rep(0,n)
25 istrain[tindex]=1
26 # Visualize
27 plot(x,col=ifelse(y>0,1,2),pch=ifelse(istrain==1,1,2))
28 legend("topleft",c('Positive Train','Positive Test','Negative Train','Negative Test'),col=c(1,1,2,2), pch=c(1,2,1,2), text.col=c(1,1,2,2))
29
30
```





```
1 data(iris)
2 attach(iris)
3
4 ## classification mode
5 # default with factor response:
6 model <- svm(Species ~ ., data = iris)
7
8 # alternatively the traditional interface:
9 x <- subset(iris, select = -Species)
10 y <- Species
11 model <- svm(x, y)
12
13 print(model)
14 summary(model)
15
16 # test with train data
17 pred <- predict(model, x)
18 # (same as:)
19 pred <- fitted(model)
20
21 # Check accuracy:
22 table(pred, y)
23
24 # compute decision values and probabilities:
25 pred <- predict(model, x, decision.values = TRUE)
26 attr(pred, "decision.values")[1:4,]
27
28 # visualize (classes by color, SV by crosses):
29 plot(cmdscale(dist(iris[, -5])),
30      col = as.integer(iris[, 5]),
31      pch = c("o", "+")[1:150 %in% model$index + 1])
32
33 ## try regression mode on two dimensions
34
35 # create data
36 x <- seq(0.1, 5, by = 0.05)
37 y <- log(x) + rnorm(x, sd = 0.2)
38
39 # estimate model and predict input values
40 m <- svm(x, y)
41 new <- predict(m, x)
42
43 # visualize
44 plot(x, y)
45 points(x, log(x), col = 2)
46 points(x, new, col = 4)
47
48 ## density-estimation
49
50 # create 2-dim. normal with rho=0:
51 X <- data.frame(a = rnorm(1000), b = rnorm(1000))
52 attach(X)
53
54 # traditional way:
55 m <- svm(X, gamma = 0.1)
56
57 # formula interface:
58 m <- svm(~., data = X, gamma = 0.1)
59 # or:
60 m <- svm(~ a + b, gamma = 0.1)
61
62 # test:
63 newdata <- data.frame(a = c(0, 4), b = c(0, 4))
64 predict (m, newdata)
65
66 # visualize:
67 plot(X, col = 1:1000 %in% m$index + 1, xlim = c(-5,5), ylim=c(-5,5))
68 points(newdata, pch = "+", col = 2, cex = 5)
69
70 # weights: (example not particularly sensible)
71 i2 <- iris
72 levels(i2$Species)[3] <- "versicolor"
73 summary(i2$Species)
74 wts <- 100 / table(i2$Species)
75 wts
76 m <- svm(Species ~ ., data = i2, class.weights = wts)
77
```

Call:
svm.default(x = x, y = y)

Parameters:
SVM-Type: C-classification
SVM-Kernel: radial
cost: 1

Number of Support Vectors: 51

Call:
svm.default(x = x, y = y)

Parameters:
SVM-Type: C-classification
SVM-Kernel: radial
cost: 1

Number of Support Vectors: 51

(8 22 21)


Number of Classes: 3

Levels:
setosa versicolor virginica


pred	y		
	setosa	versicolor	virginica
setosa	50	0	0
versicolor	0	48	2
virginica	0	2	48

A matrix: 4 × 3 of type dbl

	setosa/versicolor	setosa/virginica	versicolor/virginica
1	1.196152	1.091757	0.6708810
2	1.064621	1.056185	0.8483518
3	1.180842	1.074542	0.6439798
4	1.110699	1.053012	0.6782041



1: TRUE 2: FALSE



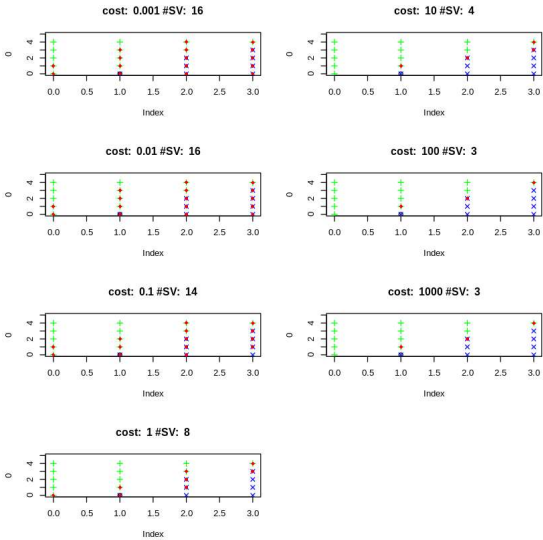
setosa: 50 versicolor: 100

setosa versicolor

2 1



```
1 library(e1071)
2 m1 <- matrix( c(
3 0,    0,    0,    1,    1,    2,    1, 2,    3,    2,    3, 3, 0, 1,2,3,
4 0, 1, 2, 3,
5 1,    2,    3,    2,    3,    3,    0, 0,    0,    1, 1, 2, 4, 4,4,4,    0,
6 1, 2, 3,
7 1,    1,    1,    1,    1,    1,    -1,-1,  -1,-1,-1,-1, 1 ,1,1,1,    1,
8 1,-1,-1
9 ), ncol = 3 )
10
11 Y = m1[,3]
12 X = m1[,1:2]
13
14 df = data.frame( X , Y )
15
16 par(mfcol=c(4,2))
17 for( cost in c( 1e-3 ,1e-2 ,1e-1, 1e0,  1e+1, 1e+2 ,1e+3)) {
18 #cost <- 1
19 model.svm <- svm( Y ~ . , data = df ,  type = "C-classification" , kernel =
20 "linear", cost = cost,
21                scale =FALSE )
22 #print(model.svm$SV)
23
24 plot(x=0,ylim=c(0,5), xlim=c(0,3),main= paste( "cost: ",cost, "#SV: ",
25 nrow(model.svm$SV) ))
26 points(m1[m1[,3]>0,1], m1[m1[,3]>0,2], pch=3, col="green")
27 points(m1[m1[,3]<0,1], m1[m1[,3]<0,2], pch=4, col="blue")
28 points(model.svm$SV[1],model.svm$SV[2], pch=18 , col = "red")
29 }
30
```



```
1 # load the kernlab package
2 install.packages("kernlab")
3 library(kernlab)
4 # train the SVM
5 svp <- ksvm(xtrain,ytrain,type="C-svc",kernel='vanilladot',C=100,scaled=c())
6 #Look and understand what svp contains
7 # General summary
8 svp
9 # Attributes that you can access
10 attributes(svp)
11 # For example, the support vectors
12 alpha(svp)
13 alphaindex(svp)
14 b(svp)
15 # Use the built-in function to pretty-plot the classifier
16 plot(svp,data=xtrain)
17
18
```

Installing package into ‘/usr/local/lib/R/site-library’
(as ‘lib’ is unspecified)

Setting default kernel parameters
Support Vector Machine object of class "ksvm"

SV type: C-svc (classification)
parameter : cost C = 100

Linear (vanilla) kernel function.

Number of Support Vectors : 3

Objective Function Value : -83.3392
Training error : 0
\$param
\$param\$C
[1] 100

\$scaling
` \001NULL\001 `

\$coef
\$coef[[1]]
[1] 78.294374 -83.301146 5.006772

\$alphaindex
\$alphaindex[[1]]
[1] 32 81 87

\$b
[1] -28.34212

\$obj
[1] -83.33923

\$SVindex
[1] 32 81 87

\$nSV
[1] 3

\$prior
\$prior[[1]]
\$prior[[1]]\$prior1
[1] 62

\$prior[[1]]\$prior0
[1] 58

\$prob.model
\$prob.model[[1]]
NULL

\$alpha
\$alpha[[1]]
[1] 78.294374 83.301146 5.006772

\$type
[1] "C-svc"

\$kernel
Linear (vanilla) kernel function.

\$kpar
list()

\$xmatrix
\$xmatrix[[1]]
X1 X2
32 1.511672 1.557370
81 1.485503 1.755448
87 -0.275778 2.658658

\$ymatrix
[1] 1 -1 1 -1 -1 -1 1 -1 -1 1 1 1 -1 1 1 1 1 -1 -1 1 -1 1 1
[26] -1 1 1 -1 -1 1 1 -1 -1 -1 1 -1 -1 -1 -1 -1 1 1 -1 -1 1 -1 1 1
[51] 1 -1 -1 -1 1 1 -1 -1 -1 -1 1 1 -1 -1 1 1 1 1 1 -1 1 1 1
[76] -1 -1 -1 -1 1 -1 1 1 -1 1 -1 1 -1 1 -1 1 1 1 -1 -1 1 1 -1 -1
[101] -1 -1 1 1 1 1 -1 1 1 1 1 -1 -1 -1 1 1 -1 1 1 -1

\$fitted
[1] 1 -1 1 -1 -1 -1 1 -1 -1 1 1 1 -1 1 1 1 1 -1 -1 1 -1 1 1
[26] -1 1 1 -1 -1 1 1 -1 -1 -1 1 -1 -1 -1 -1 -1 1 1 -1 -1 1 -1 1
[51] 1 -1 -1 -1 1 1 -1 -1 -1 1 1 -1 -1 1 1 -1 1 1 1 -1 1 1 1
[76] -1 -1 -1 -1 1 -1 1 1 -1 1 -1 1 -1 1 1 1 1 -1 -1 1 1 -1 -1
[101] -1 -1 1 1 1 1 -1 1 1 1 1 -1 -1 -1 1 1 -1 1 1 -1

\$lev
[1] -1 1

\$nclass
[1] 2

```
1 ##### Use text with string kernels
2 data(reuters)
3 is(reuters)
4 tsv <- ksvm(reuters,rlabels,kernel="stringdot",
5             kpar=list(length=5),cross=3,C=10)
6 tsv
7
8
```

'list' · 'vector' · 'listOrNULL' · 'input' · 'listl' · 'lpinut' · 'output'
Support Vector Machine object of class "ksvm"

SV type: C-svc (classification)
parameter : cost C = 10

String kernel function. Type = spectrum
Hyperparameters : sub-sequence/string length = 5
Normalized

Number of Support Vectors : 39

Objective Function Value : -13.6834
Training error : 0
Cross validation error : 0.128205

1 32 · 81 · 87